

Knowledge Management using Agentic AI in Small Scale Networks

Abstract

Traditional Knowledge Management (KM) systems are predominantly designed for enterprise-scale deployments, rendering them operationally complex, prohibitively expensive, and misaligned with the privacy constraints of small-scale networks and independent research environments. The minor phase of this project established a localized, edge-hosted document retrieval prototype utilizing a Raspberry Pi 5 to ensure strict data sovereignty. However, baseline Retrieval-Augmented Generation (RAG) pipelines utilizing naive fixed-size chunking and single-stage dense vector search exhibited severe limitations when processing complex, unstructured optical character recognition (OCR) data.

To address these deficits, this major project thesis details the architectural evolution from a passive storage utility into an autonomous, agentic knowledge repository. The core contribution is the integration of the ReAct (Reasoning and Acting) framework, empowering the system with autonomous decision-making capabilities, including web search via the Serper API and a mathematical calculator tool. Furthermore, the retrieval pipeline was fundamentally overhauled to improve context precision and factual grounding. This was achieved through the implementation of Context-Aware Metadata Chunking, a Hybrid Retrieval pipeline synthesizing dense Matryoshka embeddings with BM25 sparse lexical search, Cross-Encoder reranking, and Maximal Marginal Relevance (MMR) for cognitive diversity. System efficacy was rigorously benchmarked using the RAGAS framework. Results demonstrate a 100% Context Recall and a 94.1% Faithfulness score on the standardized SQuAD dataset, validating that enterprise-grade, agentic artificial intelligence can be executed securely and efficiently at the network edge.

Introduction

In the contemporary digital ecosystem, the proliferation of unstructured data necessitates advanced frameworks for digital archiving and retrieval. Knowledge Management encompasses the systematic processes through which entities capture, distribute, and utilize information to achieve operational and intellectual objectives. The discourse surrounding KM has historically focused on enterprise environments characterized by centralized data warehouses and vast cloud computing resources.¹ These systems prioritize rigid documentation practices serving thousands of concurrent users. However, these enterprise solutions exhibit severe misalignments when applied to small-scale networks, home laboratories, or independent research infrastructures. The financial costs of enterprise software and the privacy risks associated with delegating sensitive data to third-party cloud architectures present formidable barriers to adoption.¹

Consequently, critical information in personal environments—ranging from financial records and legal contracts to academic research—frequently becomes lost within fragmented digital silos. Traditional hierarchical folder structures and basic keyword search utilities are fundamentally misaligned with human cognitive recall mechanisms, which rely on contextual semantic cues rather than exact alphanumeric nomenclature. The absence of an intelligent retrieval system results in continuous cognitive overhead and the degradation of institutional memory.¹

To resolve these constraints, the initial minor project detailed the implementation of an edge-hosted KM system on a Raspberry Pi 5, guaranteeing absolute data sovereignty by keeping all document ingestion and vectorization strictly localized.¹ While the foundational hardware and secure networking mechanisms successfully established a zero-trust environment, the cognitive retrieval capabilities were constrained by baseline methodologies. Real-world deployment on noisy, optical character recognition extracted data revealed limitations in context preservation and generative faithfulness.¹

This major project report details the sophisticated evolution of the system. It transitions the architecture from a simple semantic search engine into an Agentic Artificial Intelligence framework. By integrating the ReAct paradigm, Hybrid Search pipelines, Cross-Encoder reranking, and Maximal Marginal Relevance, the system overcomes the chaos of unstructured physical documents, transforming static files into an interactive, multi-tool orchestrated knowledge assistant.

Literature Review

The evolution of intelligent document processing relies heavily on advancements in information retrieval and autonomous agent design. This section reviews the critical methodologies informing the major project upgrades.

Retrieval-Augmented Generation and Edge Computing RAG bridges the gap between the generalized reasoning of Large Language Models (LLMs) and domain-specific proprietary knowledge. By enforcing a deterministic retrieval step prior to generative synthesis, RAG grounds the model in empirical data, mitigating hallucinations. The minor project established that running such pipelines on edge devices like the Raspberry Pi 5 mitigates cloud dependency and latency, provided that embedding models are optimized for CPU execution, such as Matryoshka Representation Learning (MRL) architectures.¹

Hybrid Retrieval: Synthesizing Dense and Sparse Search Standard retrieval pipelines rely almost exclusively on bi-encoder dense vector search. While adept at capturing broad semantic concepts, dense retrieval frequently fails to capture exact keyword matches, acronyms, or specific numerical identifiers. Recent studies demonstrate that Hybrid Retrieval—combining semantic dense methods with sparse lexical algorithms like BM25—yields superior performance. BM25, a ranking function based on term frequency and inverse document frequency principles, efficiently handles exact keyword matching. Hybrid systems merge these paradigms, often utilizing algorithms like Reciprocal Rank Fusion (RRF) to normalize and synthesize the disparate scoring outputs.²

Cross-Encoder Reranking and Maximal Marginal Relevance While hybrid retrieval ensures high recall, it frequently sacrifices precision. Cross-Encoder reranking models address this by processing the query and the retrieved document simultaneously through transformer attention layers, capturing highly nuanced interactions that bi-encoders miss. Research from MIT indicates that adding a cross-encoder to a retrieval pipeline increases average accuracy by 33% to 40%, particularly for complex or ambiguous queries.³

To prevent the LLM from being starved of diverse context due to redundant documents, Maximal Marginal Relevance (MMR) is employed. Originally introduced by Carbonell and Goldstein (1998), MMR elegantly balances query relevance with information novelty, penalizing redundant information in the final context window to promote a richer, multi-faceted generative response.

Agentic Frameworks and ReAct

The transition from reactive AI to autonomous systems is encapsulated by agentic frameworks. The ReAct (Reasoning and Acting) paradigm, proposed by Yao et al. (2023), synergizes dynamic reasoning with task-specific actions. Unlike traditional models that only generate text, a ReAct agent iteratively formulates thoughts, executes actions via external tools (e.g., APIs, calculators), observes the results, and adjusts its plan. This tight integration allows the system to synthesize internal knowledge with dynamic, external environmental feedback.

Evaluation Frameworks: RAGAS

Evaluating RAG architectures requires specialized methodologies that do not rely strictly on human-annotated ground truth. The RAGAS (Retrieval Augmented Generation Assessment) framework, introduced by Es et al. (2023), provides a suite of metrics utilizing LLMs as objective judges. RAGAS quantitatively evaluates four primary dimensions: Faithfulness, Answer Relevancy, Context Precision, and Context Recall, offering robust diagnostic signals for pipeline optimization.


Methodology

Hardware Infrastructure and Secure Networking


The physical foundation inherited from the minor project relies on a Raspberry Pi 5 (Broadcom BCM2712 SoC, 8GB RAM) augmented with a 512GB NVMe SSD via a PCIe 2.0 interface to bypass standard microSD input/output bottlenecks.¹ To achieve "always-on" stability under the heavy thermal loads induced by continuous AI workloads, a dedicated active cooling apparatus was implemented.¹ Secure remote accessibility is maintained via Tailscale, a zero-trust mesh Virtual Private Network based on the WireGuard protocol, allowing secure ingestion and querying without exposing local ports to the public internet.¹

Baseline Pipelines


To establish a functional foundational layer, we implemented three core operational pipelines that dictated the flow of data from raw upload to final response synthesis:



1. **Automated Ingestion and OCR Pipeline:** The document digitization process begins with file entry via manual upload or an automated Samba network share.¹ The uploaded document is routed through the Tesseract Optical Character Recognition (OCR) engine. Following text extraction, the file undergoes PDF/A archival to ensure format stability. The system then performs metadata extraction and rule/pattern-based organizing before triggering a post-consumption script dedicated to vector database processing.



2. **Text Segmentation and Embedding Pipeline:** Once triggered by manual ingestion or the post-consumption script (orchestrated by Paperless-ngx), the raw text is passed into a fixed-length chunking module. As established in the minor phase, this involved rigid segments of 1000 characters with a 200-character overlap.¹ These chunks are converted into mathematical vectors via Matryoshka Representation Learning (MRL) embeddings.¹ The system subsequently creates a connection to ChromaDB to persistently store the generated embeddings alongside the raw document content and extracted metadata.



3. **Retrieval and Generation Pipeline:** The user interaction phase is initiated when a user query is submitted through the frontend interface, which sends a POST request via the FastAPI backend. The system generates an MRL embedding of the query and executes a similarity search against ChromaDB. The top k retrieved results form the context block. This context, combined with the system prompt and the original user prompt, is transmitted to the Groq API (utilizing a large 120-billion-parameter open-source model) to synthesize the final response, which is then streamed back to the frontend.¹

Context-Aware Metadata Chunking

The minor project utilized fixed-size chunking (1000 characters), which frequently severed cohesive paragraphs and detached vital section headings.[1, 1] The major project overhauled this ingestion phase to utilize Context-Aware Metadata Chunking. Advanced parsing libraries analyze the document's logical and structural hierarchy, identifying semantic breakpoints to ensure discrete concepts remain encapsulated. Furthermore, high-level metadata (document title, date, author, specific section headers) is programmatically injected directly into every individual chunk's text string prior to vectorization, permanently anchoring the fragment to its broader macro-document context.⁴

The Hybrid Retrieval Pipeline

To overcome the limitations of pure vector similarity, a comprehensive Hybrid Retrieval Pipeline was engineered:

1. **Dense Semantic Search:** Semantic embeddings are generated locally on the Raspberry Pi using a highly optimized static Matryoshka Representation Learning (MRL) model, truncated for CPU efficiency.⁵
2. **Sparse Lexical Search:** The BM25 algorithm evaluates exact lexical overlap and term-frequency statistics to capture specific alphanumeric strings and identifiers.

The outputs of both pipelines are merged utilizing Reciprocal Rank Fusion (RRF). RRF normalizes the disparate mathematical bounds of cosine similarity and BM25 scores by evaluating the ordinal rank (r_i) of a document (d) across the search lists, governed by the following formula ²:

$$Score_{RRF}(d) = \sum_{i \in \text{pipelines}} weight_i \times \frac{1}{k + r_i(d)}$$

where k is a standard smoothing constant set to 60.

Precision and Diversity Orchestration

To ensure optimal context quality for the generative model, the broad candidate pool retrieved via Hybrid Search undergoes a strict, two-stage filtration process:

1. **Cross-Encoder Reranking:** A secondary evaluation is executed using a Cross-Encoder model. The query and each candidate chunk are passed simultaneously through the model to map intricate contextual dependencies. Irrelevant chunks retrieved via tangential semantic drift are aggressively down-scored and discarded, functioning as a strict judicial filter.
2. **Maximal Marginal Relevance (MMR):** To enforce cognitive diversity and prevent context window saturation with redundant boilerplate text, the reranked list is processed via an MMR algorithm. The selection balances strict query relevance and informational novelty using the penalty function:

$$MMR = \lambda \times \text{Relevance} - (1 - \lambda) \times \text{Similarity}$$

This algorithm iteratively evaluates the pairwise cosine similarity of new candidates against already-selected documents, ensuring a diverse spread of factual data.⁶

ReAct Agentic Framework and Tool Orchestration

The most profound architectural upgrade is the transition from a passive retrieval script to the ReAct agentic framework. Operating as an autonomous knowledge worker, the agent parses user inputs and engages in an iterative Thought -> Action -> Observation loop.

When a query is submitted, the agent evaluates if internal knowledge retrieval is sufficient. If the query demands external real-time data or complex mathematical deduction, the agent autonomously invokes distinct tools:

- **Web Search Integration:** Integrated via the Serper API, this tool allows the agent to execute real-time, semantic Google searches to fetch contemporary information, answering queries that fall beyond the temporal cutoff of the localized database.
- **Mathematical Calculator:** A dedicated calculator tool enables the agent to evaluate mathematical expressions accurately, mitigating the inherent arithmetic weaknesses of standard LLMs.

The orchestration relies on GroqCloud's high-velocity inference Application Programming Interface (API) to execute the LLM reasoning loop rapidly, while the edge device exclusively manages the privacy-critical vector management, document ingestion, and tool execution.¹

Result and Analysis


To rigorously validate the efficacy of the advanced architectural upgrades (Hybrid Search, Cross-Encoders, MMR, and the ReAct framework), the system underwent exhaustive quantitative benchmarking utilizing the RAGAS framework.⁷ Evaluation focused on Faithfulness (factual grounding), Answer Relevancy, Context Precision, and Context Recall.⁷

The advanced pipeline was benchmarked against the standardized Stanford Question Answering Dataset (SQuAD) and a proprietary synthetic database comprising the system's actual real-world data (scanned receipts, noisy OCR extractions, and fragmented text).⁸

Metric	SQuAD (100 questions)	Your DB (20 synthetic)	Improvement
Faithfulness	94.1%	77.5%	+16.6%
Answer Relevancy	85.1%	50.8%	+34.3%
Context Precision	81.7%	52.3%	+29.4%
Context Recall	100%	63.8%	+36.2%

Analysis of SQuAD Performance:

The system demonstrated exceptional proficiency on the SQuAD dataset. The perfect Context Recall (100%) unequivocally validates the implementation of the Hybrid Search pipeline; the combination of BM25 and dense Matryoshka vectors ensured that the necessary ground-truth information was retrieved in every single test instance. Furthermore, the high Context Precision (81.7%) proves the efficacy of the Cross-Encoder reranking module, successfully pushing the most critical information to the top of the context window. This pristine data curation directly resulted in a Faithfulness score of 94.1% and Strong Relevancy (85.1%), indicating that hallucinations were virtually eliminated when processing well-structured data.



Analysis of Synthetic Database Performance:

The performance divergence observed in the synthetic database highlights the inherent complexities of edge-based document management. SQuAD contexts consist of clean, highly structured Wikipedia paragraphs (200-400 words), representing an ideal environment for vector retrieval. In contrast, the synthetic database consists of complex, fragmented OCR text derived from physical documents.

The lower scores across all metrics (Context Recall dropping to 63.8% and Answer Relevancy to 50.8%) suggest that the original physical documents are highly fragmented, making it difficult for the parsing engines to maintain semantic continuity. However, sustaining a 77.5% Faithfulness score on severely degraded data proves the absolute necessity of the implemented agentic upgrades. Without the Cross-Encoder filtering out noisy OCR errors and the MMR algorithm ensuring diverse context mapping, real-world performance on such chaotic data would degrade to unusable levels. The RAG pipeline works exceptionally well with clean documents, but physical edge deployments must account for OCR-induced entropy.

Conclusion and Future Work

This major project successfully transitions a baseline, passive document storage prototype into a highly autonomous, agentic Knowledge Management system optimized for secure, edge-network execution. By overhauling the underlying retrieval mechanics to include Context-Aware Chunking, Hybrid Dense and Sparse Search, and advanced Cross-Encoder filtering, the architecture effectively overcomes the severe limitations of standard vector databases when applied to noisy real-world data. Furthermore, the integration of the ReAct paradigm, complete with external Serper API web search and calculator tools, transforms the system into a dynamic reasoning engine capable of complex, multi-hop problem-solving. Achieving perfect retrieval recall and 94.1% faithfulness on standardized benchmarks confirms that enterprise-grade AI capabilities can be democratized and executed on affordable, low-power hardware like the Raspberry Pi 5 without compromising data sovereignty.

Future Work:

While the current ingestion pipeline relies heavily on traditional OCR engines, extracting text from complex visual layouts (financial tables, academic graphs) inevitably destroys spatial relationships. Future iterations of this architecture will aim to deprecate traditional OCR entirely in favor of lightweight, edge-compatible Vision-Language Models (VLMs). VLMs process document pages natively as images, capturing the exact spatial relationships of tables and figures, fundamentally bridging the gap between human visual comprehension and machine indexing. Additionally, future work will focus on integrating persistent Agentic Memory structures, allowing the system to archive multi-turn dialogue trajectories to facilitate long-horizon conversational continuity.

References

1. Es, S., James, J., Espinosa-Anke, L., & Schockaert, S. (2023). Ragas: Automated Evaluation of Retrieval Augmented Generation. *arXiv preprint arXiv:2309.15217*.
2. Yao, S., Zhao, J., Yu, D., Du, N., Shafran, I., Narasimhan, K., & Cao, Y. (2023). ReAct: Synergizing Reasoning and Acting in Language Models. *International Conference on Learning Representations (ICLR)*.
3. Carbonell, J. G., & Goldstein, J. (1998). The Use of MMR, Diversity-Based Reranking for Reordering Documents and Producing Summaries. *Proceedings of the 21st Annual International ACM SIGIR Conference on Research and Development in Information Retrieval*.
4. Robertson, S. E., Walker, S., Jones, S., Hancock-Beaulieu, M. M., & Gatford, M. (1994). Okapi at TREC-3. *Overview of the Third Text REtrieval Conference (TREC-3)*.
5. Nogueira, R., & Cho, K. (2020). Passage Re-ranking with BERT. *arXiv preprint arXiv:1901.04085*.
6. Nguyen, L. T. K., Connolly, J., & Nguyen, H. N. (2025). A Systematic Review of Improving Knowledge Management with Generative AI and Large Language Models. *Journal of Advances in Information Technology*, 16(4), 594–601.
7. Lee, C., Lee, G., & Lin, H. (2007). The role of organizational capabilities in successful e-business implementation. *Business Process Management Journal*, 13(5), 677–693.
8. Rzhask, S. (2025). Leveraging Artificial Intelligence for Knowledge Management: A Systematic Literature Analysis.
9. Al-Shamaileh, K. M., & Al-Shamaileh, M. A. (2022). Integrating artificial intelligence for knowledge management systems – synergy among people and technology. *Cogent Business & Management*, 9(1).
10. Abrizah, A., Zainab, A. N., Kiran, K., & Raj, R. G. (2013). LIS journals scientific impact and subject categorization: a systematic review. *Scientometrics*.
11. Al-Emran, A. (2024). AI-Based Knowledge Management Systems: A Review of AI Techniques, Applications and Challenges.
12. Anshari, M., Syafrudin, M., Tan, A., Fitriyani, N. L., & Alas, Y. (2023). Optimisation of Knowledge Management with Machine Learning Enabled. *Information*, 14(1), 35.
13. Reimers, N., & Gurevych, I. (2019). Sentence-BERT: Sentence Embeddings using Siamese BERT-Networks. *Proceedings of EMNLP*.
14. Lewis, P. et al. (2020). Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks. *Advances in Neural Information Processing Systems (NeurIPS)*, 33, 9459–9474.
15. Izacard, G., & Grave, E. (2021). Leveraging Passage Retrieval with Generative Models for Open-Domain Question Answering. *Proceedings of EACL*.
16. Karpukhin, V. et al. (2020). Dense Passage Retrieval for Open-Domain Question Answering. *Proceedings of EMNLP*.
17. Harish, H. M. (2024). The Intelligent Edge: AI and Machine Learning in Edge Computing for IoT.
18. Khan, S. A. (2024). Edge AI: Deploying Artificial Intelligence Models on Edge Devices for Real-Time Analytics.
19. Shi, W., Cao, J., Zhang, Q., Y. Li, & Xu, L. (2016). Edge Computing: Vision and Challenges. *IEEE Internet of Things Journal*, 3(5), 637–646.

21. Zhang, C., & Dabak, A. (2021). Edge AI: On-Device Machine Learning for Internet of Things Applications. *IEEE Internet of Things Magazine*, 4(3), 6–12.
22. Malkov, Y. A., & Yashunin, D. A. (2020). Efficient and Robust Approximate Nearest Neighbor Search Using Hierarchical Navigable Small World Graphs. *IEEE Transactions on Pattern Analysis and Machine Intelligence*, 42(4), 824–836.
23. Johnson, J., Douze, M., & Jégou, H. (2021). Billion-Scale Similarity Search with GPUs. *IEEE Transactions on Big Data*, 7(3), 535–547.
24. Kusupati, A. et al. (2022). Matryoshka Representation Learning. *Advances in Neural Information Processing Systems (NeurIPS)*.
25. Singh, S. K., Singh, A. K., & Singh, R. K. (2025). Agentic AI: A Comprehensive Survey of Technologies, Applications, and Societal Implications. *IEEE Access*.
26. Al-Ani, A. T. (2025). AI Agents vs. Agentic AI: A Conceptual Taxonomy, Applications and Challenges. *arXiv preprint arXiv:2505.10468*, 2025.